

# Multi-scale CAFE framework for simulating fracture in heterogeneous materials implemented in Fortran coarrays and MPI

A. Shterenlikht<sup>1</sup>, L. Margetts<sup>2</sup>, L. Cebamanos<sup>3</sup>,  
J.D. Arregui-Mena<sup>2</sup>

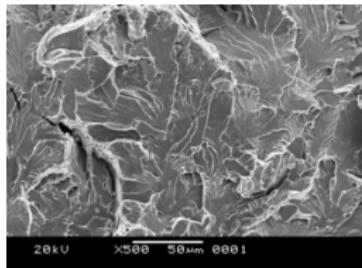
<sup>1</sup>Mech Eng Dept, The University of Bristol  
Bristol BS8 1TR, UK, Email: [mexas@bris.ac.uk](mailto:mexas@bris.ac.uk)

<sup>2</sup>School of Mechanical, Aero and Civil Engineering, The University of Manchester,  
Manchester M13 9PL, UK, Email: [Lee.Margetts@manchester.ac.uk](mailto:Lee.Margetts@manchester.ac.uk)

<sup>3</sup>Edinburgh Parallel Computing Centre (EPCC), The University of Edinburgh,  
King's Buildings, Edinburgh EH9 3FD, UK, Email: [l.cebamanos@epcc.ed.ac.uk](mailto:l.cebamanos@epcc.ed.ac.uk)

PGAS Applications Workshop, MON 14-NOV-2016  
Held in conjunction with SC16

# Fracture in heterogeneous materials



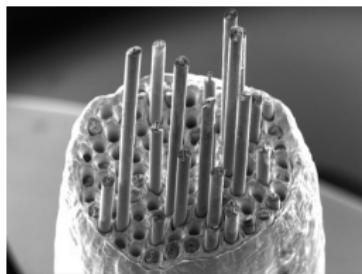
polycrystal cleavage



reinforced concrete



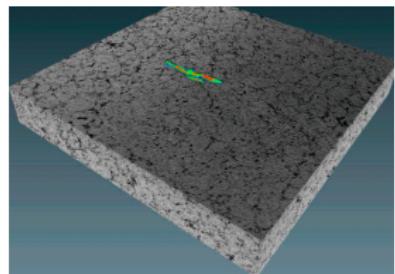
CFRP



metal matrix



bone



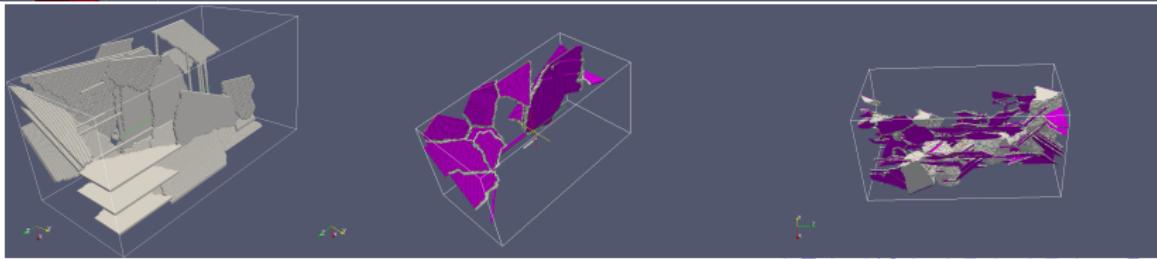
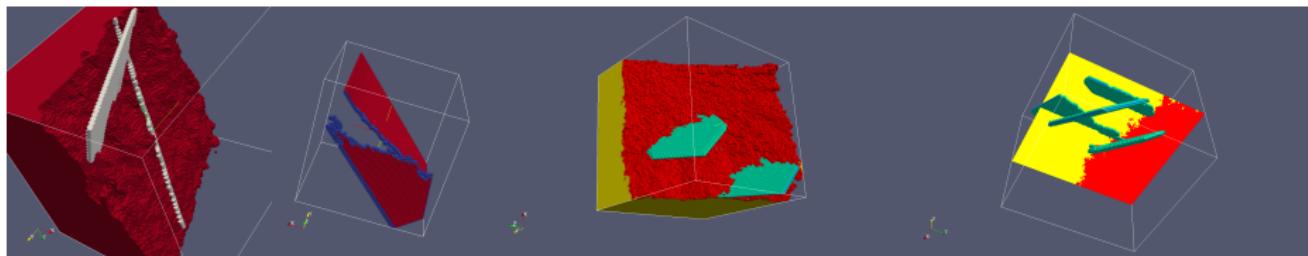
graphite

- ▶ All real materials are heterogeneous
- ▶ Multiple fracture and damage processes happen at different time and length scales → **need multi-scale framework**

# Fracture: CA + FE = CAFE multi-scale model

▶ CGPACK

- ▶ Structured grids - cellular automata (CA), unstructured grids - finite elements (FE)
- ▶ CA (microstructure) + FE (continuum mechanics) = CAFE
- ▶ Transgranular cleavage - fracture stress or strain criteria
- ▶ FE → CA (localisation) - stress, strain fields
- ▶ CA → FE (homogenisation) - damage variables

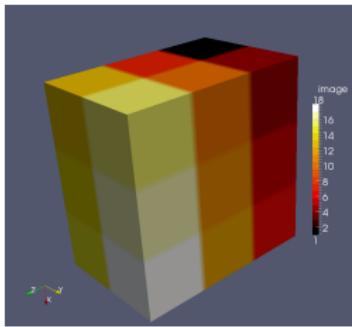


# Fortran coarrays for CA

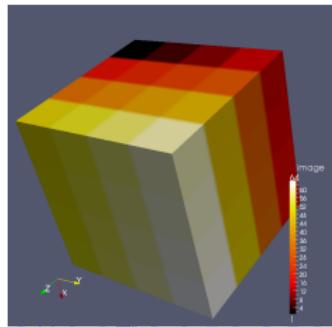
- ▶ Fortran native SPMD parallel programming feature
- ▶ Fortran standard since 2008. More features in 2015.
- ▶ Cray, Intel, OpenCoarrays/GCC support
- ▶ CGPACK - cellular automata microstructure simulation library: [cgpack.sf.net](http://cgpack.sf.net). See also [1, 2, 3].
- ▶ Easy halo exchange
- ▶ CA space coarray - 4D array, 3 codimensions:

```
integer , allocatable :: space (:,:,:,:) [:,:,:]
```

- ▶ Ideal for **structured** grids:

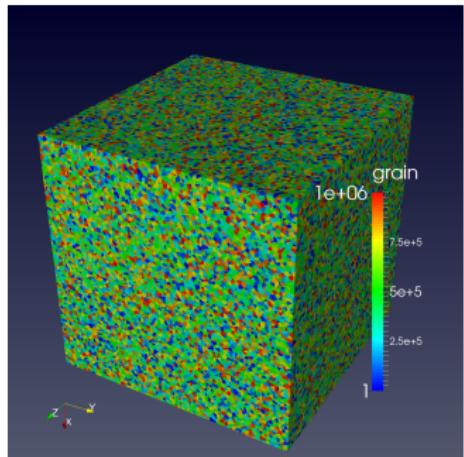
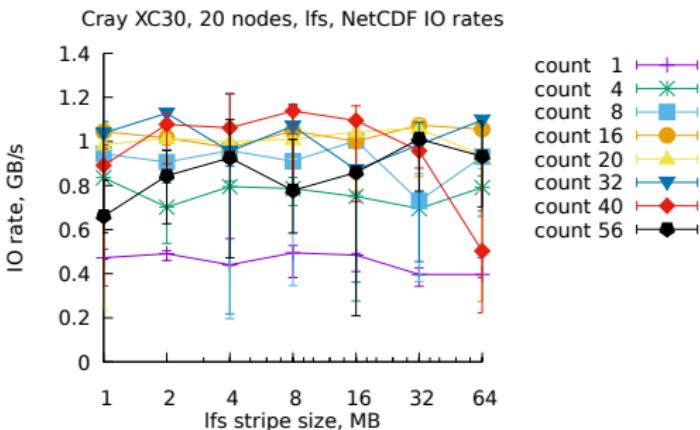


← 18 imgs; 64 imgs →

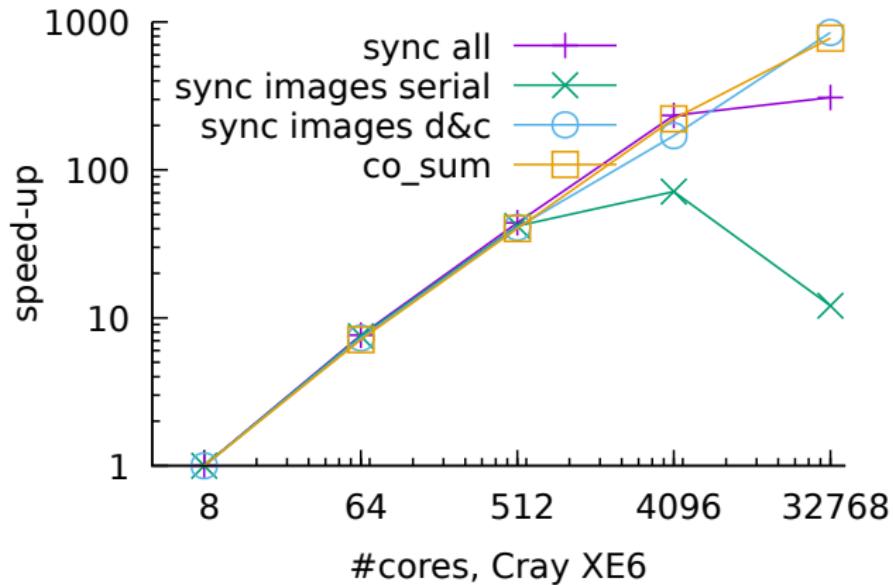


# Coarray IO - no native Fortran parallel IO

- ▶ MPI/IO up to 2.3GB/s on Cray XE6 [BCS talk](#)
- ▶ MPI/IO up to 8GB/s on Cray XC30 (can reach 14GB/s [4])
- ▶ NetCDF 4.3, HDF5 1.8.14 - only up to 1.2GB/s on Cray XC30.
- ▶ Ifs stripe count, size, number of images, file size, Cray hugepages...
- ▶ 0.5 - 1TB datasets



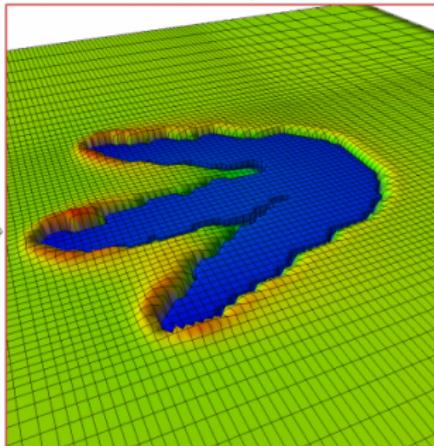
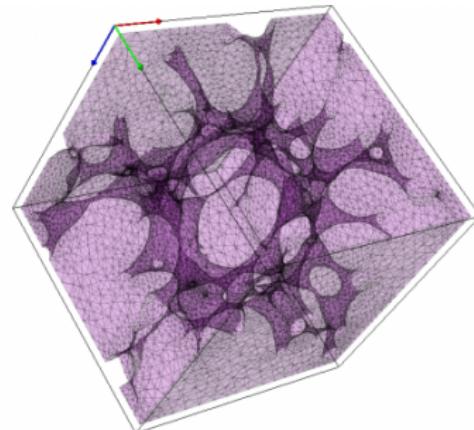
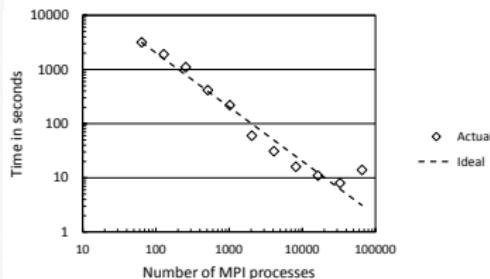
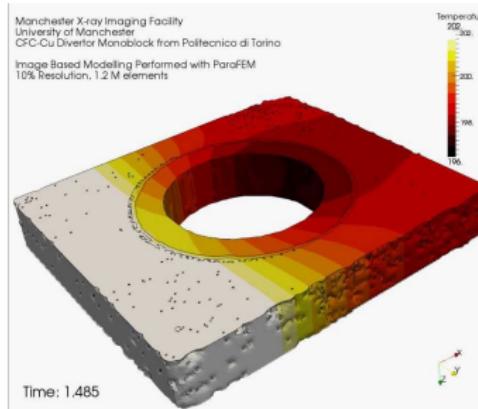
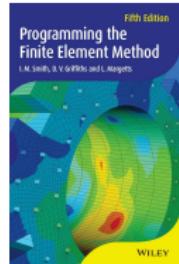
## CGPACK solidification scaling



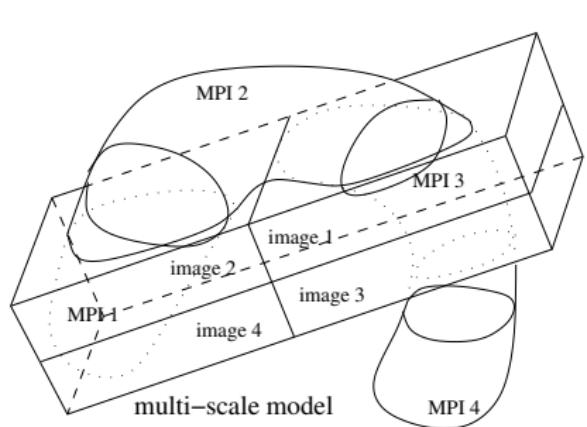
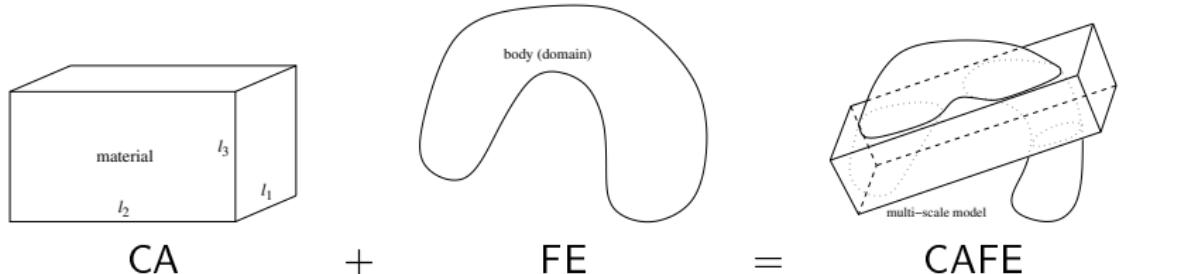
Scaling varies for different programs built with CGPACK, depending on which routines are called, in what order and requirements for synchronisation.

# ParaFEM - scalable general purpose finite element library

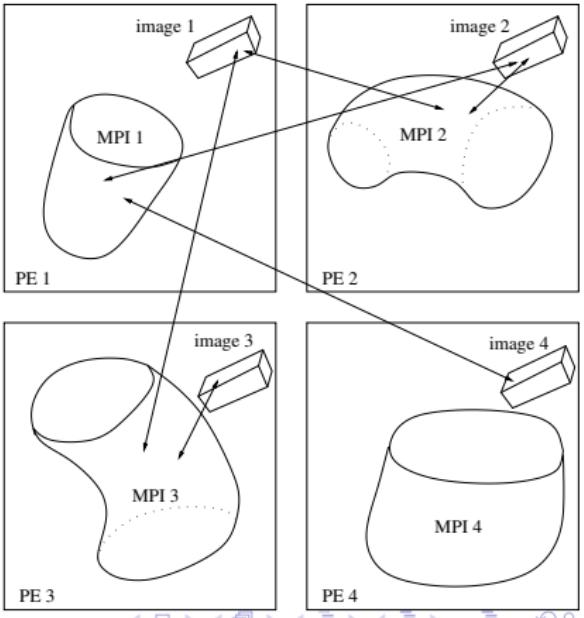
- ▶ <http://parafem.org.uk>
- ▶ Fortran 90 MPI
- ▶ Highly portable, many users [5]
- ▶ Excellent scaling
- ▶ BSD license



# CAFE design: structured CA grid + unstructured FE grid



Example with 4 PE (4 MPI processes, 4 coarray images). Arrows are  $\text{FE} \leftrightarrow \text{CA}$  comms.



FE → CA mapping via a private allocatable array of derived type:

```
type mcen
    integer :: image, elnum
    real :: centr(3)
end type mcen
type( mcen ), allocatable :: lcentr(:)
```

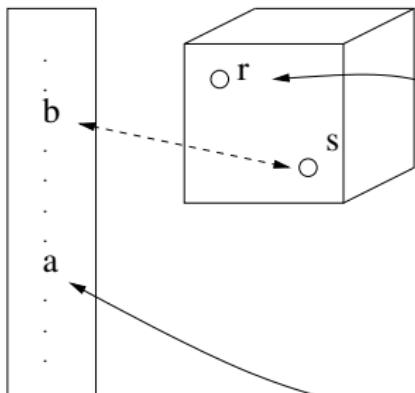
based on coordinates of FE centroids calculated by each MPI process and stored in centroid\_tmp coarray:

```
type rca
    real, allocatable :: r(:,:,:)
end type rca
type( rca ) :: centroid_tmp [*]
:
allocate( centroid_tmp%r(3, nels_pp) )
```

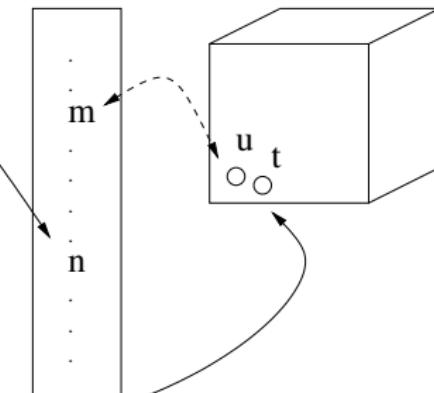
where nels\_pp is the number of FE stored on this PE.

# lcentr arrays on images P and Q

PE, image, MPI process P  
elements CA



PE, image, MPI process Q  
elements CA



lcentr

lcentr

image  
elnum  
centr

...	Q	...	P	...
...	n	...	b	...
...	r	...	s	...

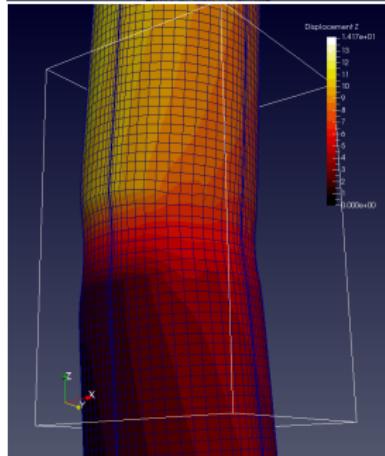
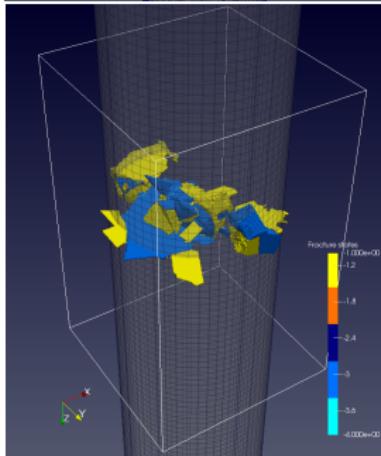
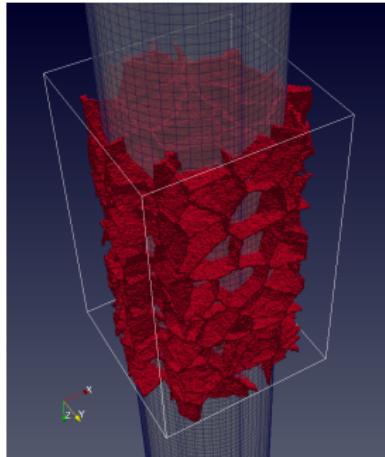
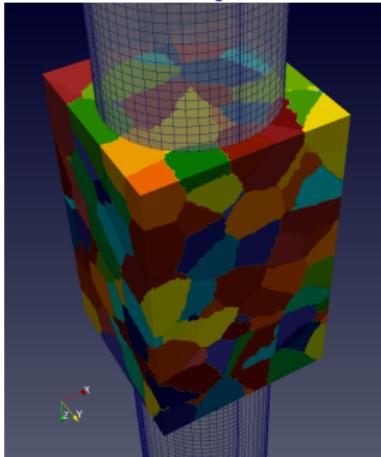
image  
elnum  
centr

...	Q	...	P	...
...	m	...	a	...
...	u	...	t	...

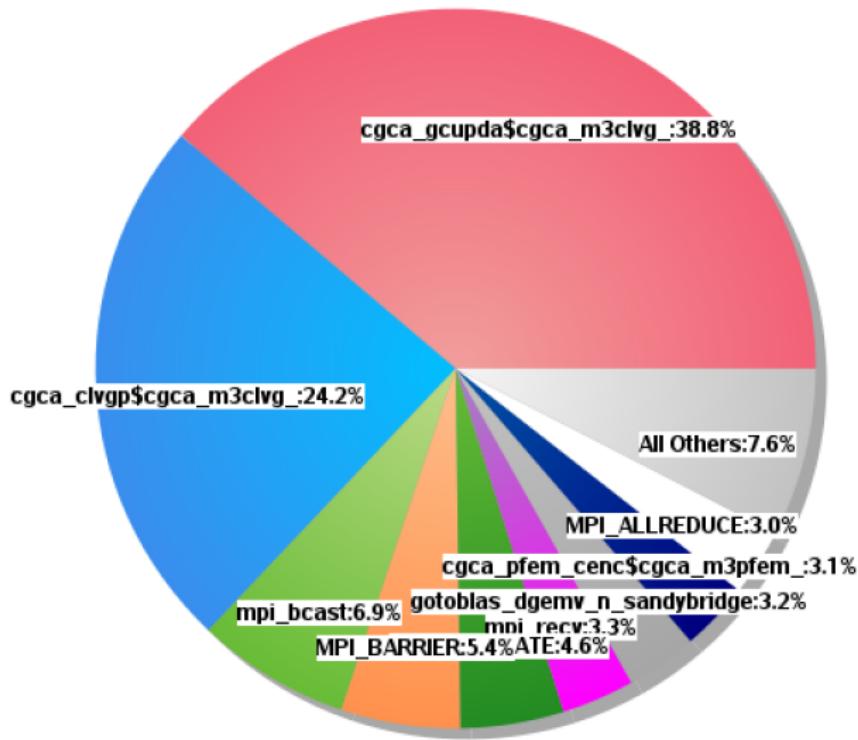
## Fracture modelling

- ▶ Diverse CAFE fracture models can be constructed from CGPACK + ParaFEM libraries.
- ▶ Simple case: isotropic linear elastic FE ( $E, \nu$ ) + cleavage (fully brittle transgranular fracture mode) CA.
- ▶ FE stress tensor  $\mathbf{t}$  passed to CA, resolved on normal stresses on  $\{100\}$  and  $\{110\}$  crystal planes -  $t_{100}, t_{110}$  [1, 6].
- ▶ 2 parameters - fracture stress,  $\sigma_F$ , linked to the free surface energy,  $\gamma$ , and a characteristic length,  $L$ .
- ▶ If  $t_{100} \geq \sigma_F$  or  $t_{110} \geq \sigma_F$  then a CA crack extends by  $L$  per unit of time.
- ▶ Crack morphology is reduced to a single damage variable,  $d$ .  $d = 1$  initially (no damage).  $d = 0$  - integration point has failed, no load bearing capacity.

# Cleavage in a steel cylinder under tension



# CrayPAT profiling 1



Profiling function distribution for ParaFEM/CGPACK MPI/coarray miniapp with all-to-all routine cgca\_gcupda at 7200 cores.

# CrayPAT profiling 2

100.0%	20,520.4	--	--	Total
-----				
71.4%	14,649.9	--	--	USER
-----				
38.7%	7,950.6	913.4	10.3%	cgca_gcupda\$cgca_m3clvg_
24.1%	4,951.2	940.8	16.0%	cgca_clvgp\$cgca_m3clvg_
3.1%	638.0	70.0	9.9%	cgca_pfem_cenc\$cgca_m3pfem_
1.8%	367.5	578.5	61.2%	cgca_hxi\$cgca_m2hx_
1.7%	346.0	196.0	36.2%	cgca_clvgn\$cgca_m3clvg_
=====				
19.8%	4,061.4	--	--	MPI
-----				
6.9%	1,413.5	356.5	20.1%	mpi_bcast
5.4%	1,098.3	419.7	27.7%	MPI_BARRIER
3.3%	670.0	322.0	32.5%	mpi_recv
3.0%	615.3	61.7	9.1%	MPI_ALLREDUCE
=====				
8.8%	1,797.2	--	--	ETC
-----				
4.6%	950.5	5.5	0.6%	__DEALLOCATE
3.2%	654.2	110.8	14.5%	gotoblas_dgemv_n_sandybridge
=====				

Raw profiling data for ParaFEM/CGPACK MPI/coarray miniapp  
with all-to-all routine cgca\_gcupda at 7200 cores.

## cgca\_gcupda - all-to-all

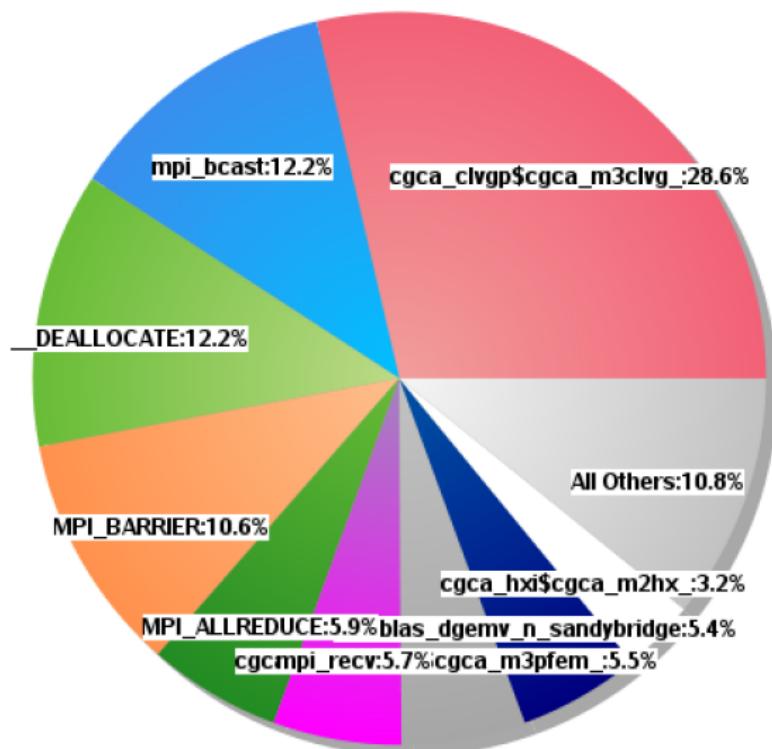
```
integer :: gcupd(100,3)[*], rndint, j, &
           img, gcupd_local(100,3)
real :: rnd
:
call random_number( rnd )
rndint = int( rnd*num_images() ) + 1
do j = rndint, rndint + num_images() - 1
  img = j
  if (img .gt. num_images()) &
      img = img - num_images()
  if (img .eq. this_image()) cycle
  :
  gcupd_local(:, :) = gcupd(:, :)[:, img]
  :
end do
```

## cgca\_gcupdn - nearest neighbour

```
do i = -1 , 1
do j = -1 , 1
do k = -1 , 1
! Get the coindex set of the neighbour
ncod = mycod + (/ i, j, k /)
:
gcupd_local(:, :) = &
    gcupd(:, :)[ncod(1), ncod(2), ncod(3)]
:
end do
end do
end do
```

Note: the nearest neighbour must be called *multiple times* to propagate changes from every image to all other images.

## CrayPAT profiling cgca\_gcupdn



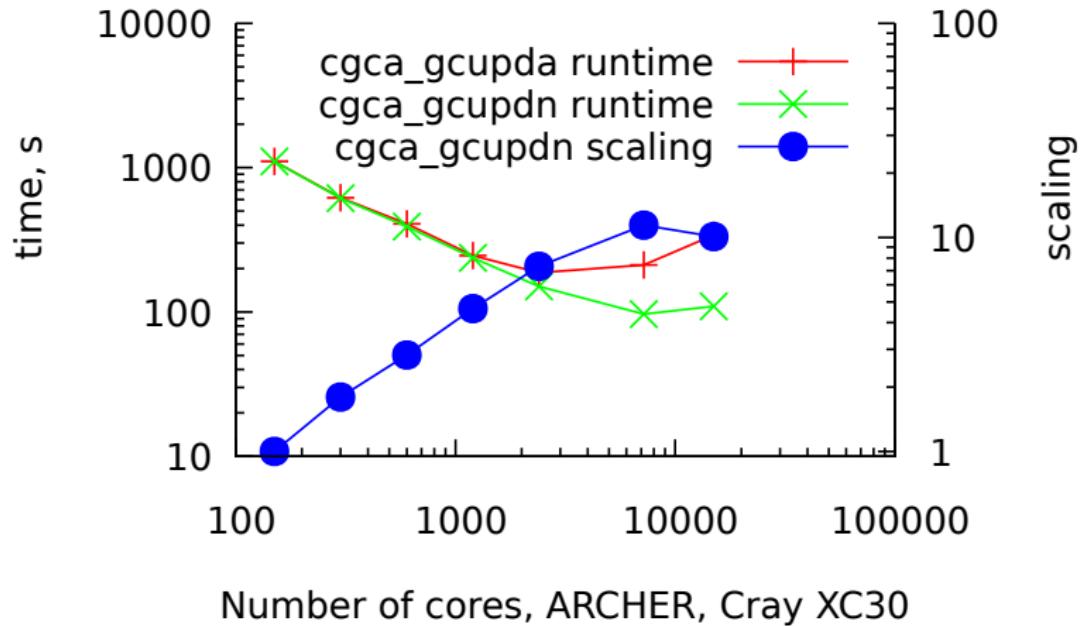
Profiling function distribution for ParaFEM/CGPACK MPI/coarray miniapp with the nearest neighbour routine cgca\_gcupdn at 7200 cores.

# CrayPAT profiling cgca\_gcupdn

100.0%	12,199.5	--	--	Total
-----				
44.8%	5,459.7	--	--	USER
-----				
28.6%	3,484.0	582.0	14.3%	cgca_clvgp\$cgca_m3clvg_
5.5%	666.1	93.9	12.4%	cgca_pfem_cenc\$cgca_m3pfem_
3.2%	393.1	752.9	65.7%	cgca_hxi\$cgca_m2hx_
2.8%	346.0	176.0	33.7%	cgca_clvgn\$cgca_m3clvg_
1.4%	165.2	37.8	18.6%	cgca_sld\$cgca_m3sld_
1.0%	126.0	82.0	39.4%	xx14_
=====				
36.7%	4,472.1	--	--	MPI
-----				
12.2%	1,484.4	380.6	20.4%	mpi_bcast
10.6%	1,287.9	389.1	23.2%	MPI_BARRIER
5.9%	714.9	90.1	11.2%	MPI_ALLREDUCE
5.7%	689.4	338.6	32.9%	mpi_recv
1.5%	179.1	417.9	70.0%	MPI_REDUCE
=====				
18.5%	2,256.1	--	--	ETC
-----				
12.1%	1,480.9	4.1	0.3%	__DEALLOCATE
5.4%	653.8	95.2	12.7%	gotoblas_dgemv_n_sandybridge
=====				

Raw profiling data for ParaFEM/CGPACK MPI/coarray miniapp  
with the nearest neighbour routine cgca\_gcupdn at 7200 cores.

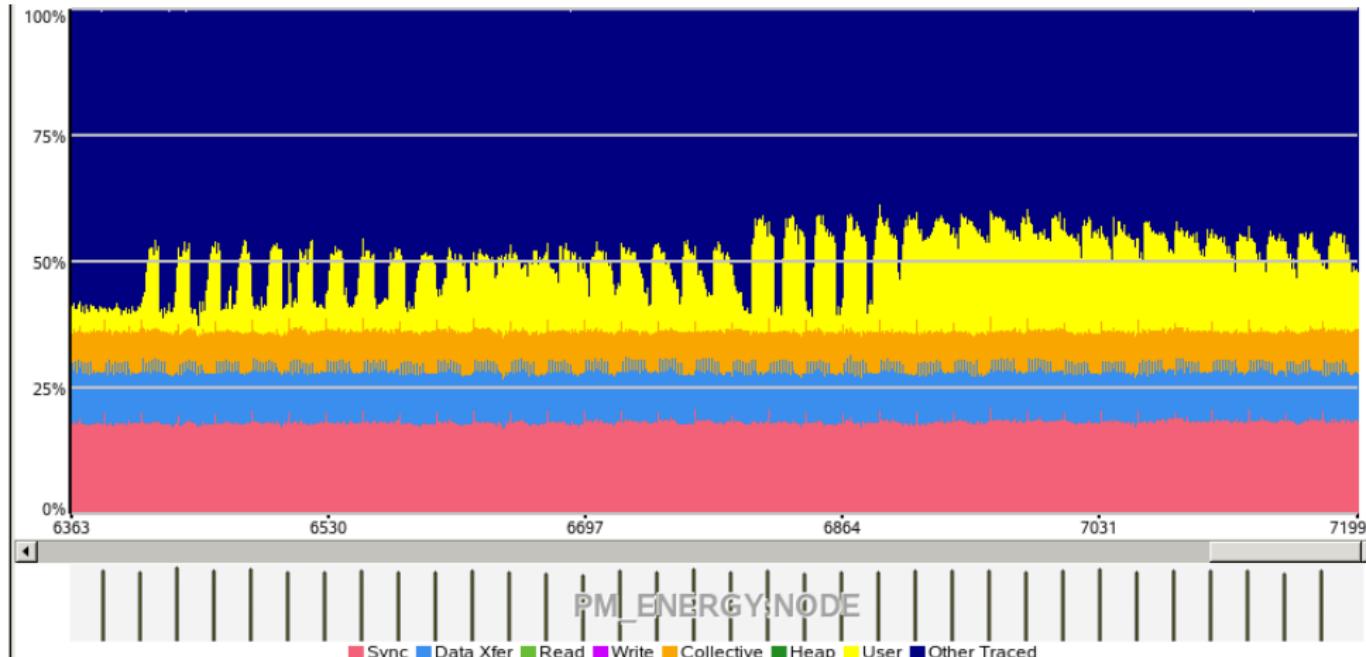
## Scaling improvement with cgca\_gcupdn over cgca\_gcupda



Number of cores, ARCHER, Cray XC30

Runtimes and scaling for ParaFEM/CGPACK MPI/coarray miniapp with the nearest neighbour, cgca\_gcupdn, and all-to-all, cgca\_gcupda, algorithms.  
Scaling limit increased from 2k to 7k cores.

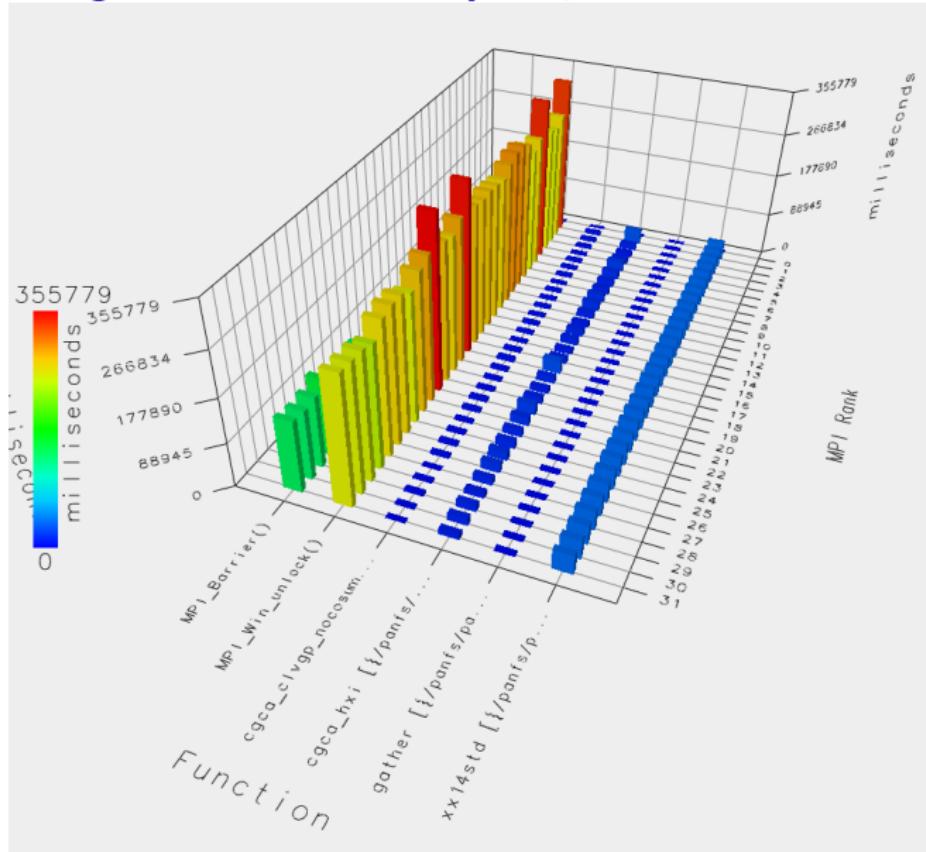
# CrayPAT - load imbalance on 7k cores on Cray XC30



Whole program activity, shown in % total time per process.

Processes 6363 to 7199 are shown.

TAU profiling: Intel 16 coarray implementation - MPI RMA



2x 16-core nodes, 32 images. Poor optimisation?

## CA - coarray (over)synchronisation?

```
call cgca_nr( space ) ! sync all inside
call cgca_rt( grt ) ! sync all inside
call cgca_sld( space ) ! sync all inside
call cgca_igb( space )
sync all
call cgca_hxi( space )
sync all
call cgca_gbs( space )
sync all
call cgca_hxi( space )
sync all
call cgca_gcu( space ) ! local routine no sync
```

- ▶ All images sync with their 26 neighbours.
- ▶ Some routines have sync inside.
- ▶ Other sync responsibility is left to end user.

## Fortran 2015 events: more flexible than SYNC IMAGES [7]

```
use, intrinsic iso_fortran_env, only:event_type
type(event_type) :: var[:, :, :]
integer, allocatable :: space(:, :, :, :, :)[:, :, :]
integer :: errstat, myrank(3)
! allocate var, space
myrank = this_image( space )
! do some work, then notify neighbours
event post(var[myrank(1)-1,myrank(2),myrank(3)], &
    stat=errstat)
! 25 more posts
:
event wait(var, until_count=26, stat=errstat)
! when all 26 neighbours posted, continue work
:
```

## Future: thread level parallelism: OpenMP, DO CONCURRENT

```
main: do iter = 1,N
do x3 = lbr(3), ubr(3)
do x2 = lbr(2), ubr(2)
do x1 = lbr(1), ubr(1)
  live: if ...
    call cgca_clvgn( clvgflag )
    if ( clvgflag ) call sub( space )
  end if live
end do
end do
end do
call co_sum( clvgglob )
sync all
call cgca_hxi( space )
sync all
call cgca_dacf( space )
```

# Conclusions

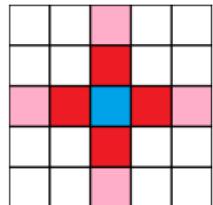
- ▶ Fortran coarrays are an ideal match for cellular automata
- ▶ Hybrid coarray+MPI multi-scale fracture framework is feasible
- ▶ Scaling up to 7k cores currently, work ongoing
- ▶ Profiling/tracing tools: CrayPAT, TAU, Score-P, Scalasca - coarray support is improving
- ▶ Coarray synchronisation - major issue: data integrity & performance

# Acknowledgements

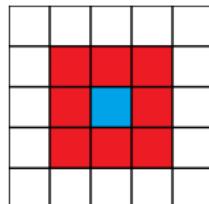
- ▶ This work was funded under the embedded CSE programme of the ARCHER UK National Supercomputing Service (<http://www.archer.ac.uk>), [► archer.ac.uk](http://www.archer.ac.uk)
- ▶ This work was carried out using the computational facilities of the Advanced Computing Research Centre, University of Bristol - <https://www.acrc.bris.ac.uk>, [► www.acrc.bris.ac.uk](https://www.acrc.bris.ac.uk)

# Cellular automata (CA) basics

- discrete space, discrete time, discrete states - fully digital framework, **structured grids**
- finite or infinite space
- finite space: fixed or self-similar boundaries
- cell neighbourhood, e.g. von Neumann's:



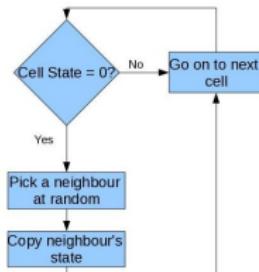
or Moore's:



- iterative process
- state of a cells at next iteration is a function of the state of this cells and of the states of its neighbourhood cells at the current iteration

# Primitive 3D solidification - probabilistic CA

- ▶ States: liquid = 0, crystals > 0.
- ▶ Cell state uniquely encodes crystal orientation tensor, i.e. a look-up table.
- ▶ Each iteration a liquid cell acquires a state of a randomly chosen neighbour (3D Moore's neighbourhood - 26 cells).



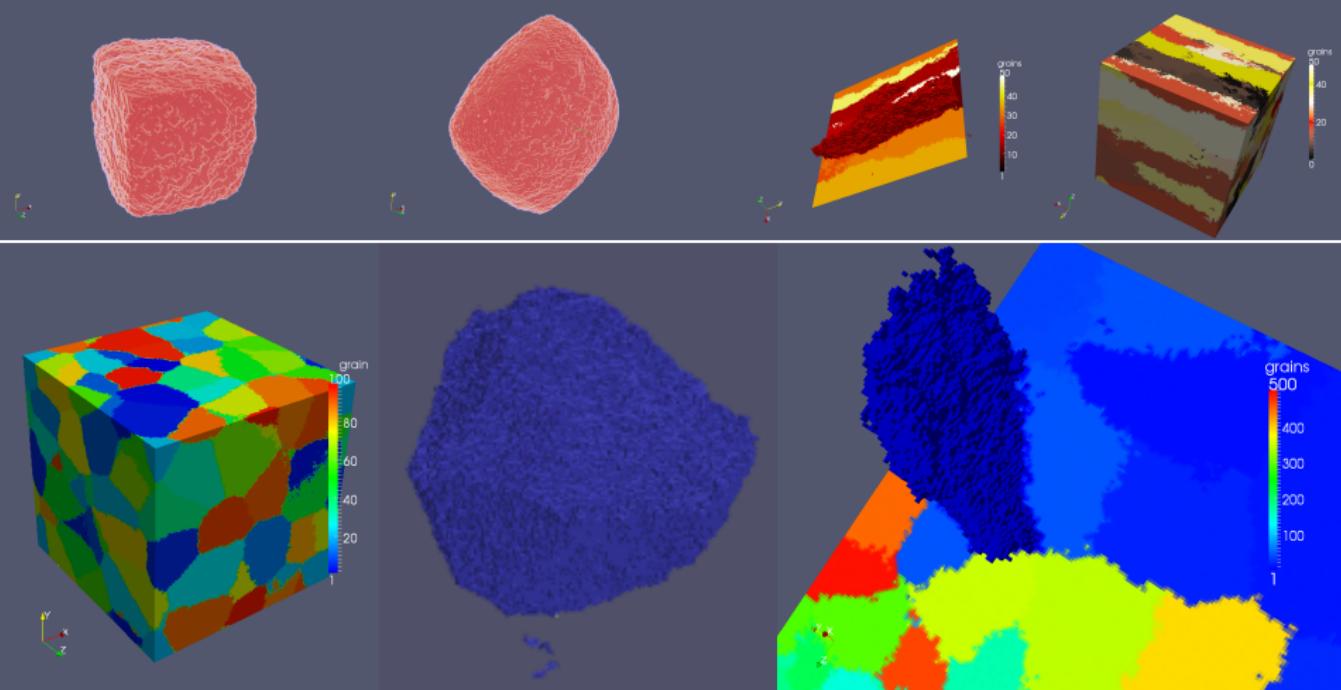
0	0	0	0	0	0	1	1	1	1	1
0	0	0	0	0	1	1	1	1	1	1
0	0	0	0	0	1	1	1	1	1	1
0	0	0	0	0	0	1	1	1	24	24
0	0	0	0	0	0	24	1	24	24	
0	0	0	0	0	0	24	24	24	24	
0	0	0	0	0	0	24	24	24	24	24
0	0	0	0	0	24	24	24	24	24	24
0	0	0	0	24	24	24	24	24	24	24
0	0	0	24	24	24	24	24	24	24	24
0	0	24	24	24	24	24	24	24	24	24

*i*

0	0	0	1	1	1	1	1	1	1
0	0	0	0	1	1	1	1	1	1
0	0	0	1	1	1	1	1	1	1
0	0	0	0	24	1	1	24	24	
0	0	0	0	0	24	1	24	24	
0	0	0	0	0	0	24	24	24	24
0	0	0	0	0	0	24	24	24	24
0	0	0	24	24	24	24	24	24	24
0	0	24	24	24	24	24	24	24	24
0	0	24	24	24	24	24	24	24	24

*i + 1*

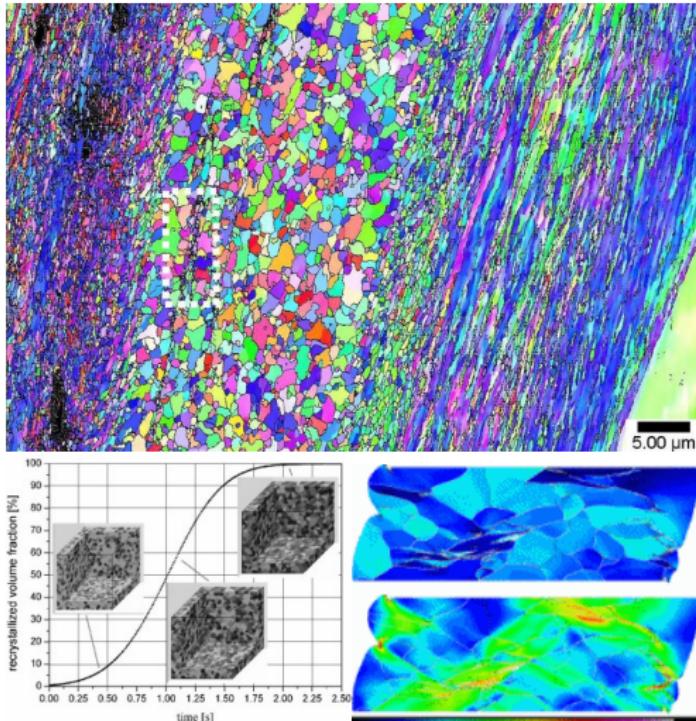
# Primitive probabilistic 3D solidification - results



For more results ➔ CGPACK

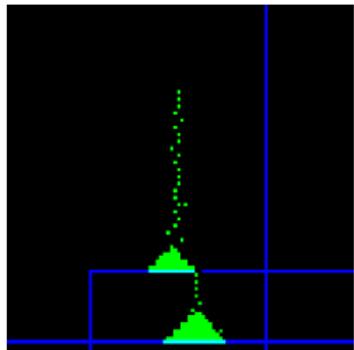
# Recrystallisation

- ▶ The grain boundary velocity  $\dot{\mathbf{x}} = \mathbf{n}mp$ ,  $\mathbf{n}$  - the normal to the grain boundary segment,  $m$  - mobility,  $p$  - the driving force.
- ▶ If  $\dot{\mathbf{x}}\Delta t \geq c$ , where  $\Delta t$  - time increment,  $c$  - cell size, then a cell joins the growing grain.
- ▶ Mobility strongly depends on temperature:  
 $m \approx \alpha \exp(-\beta/T)$ ,  $\alpha, \beta$  - some parameters,  $T$  - temperature.

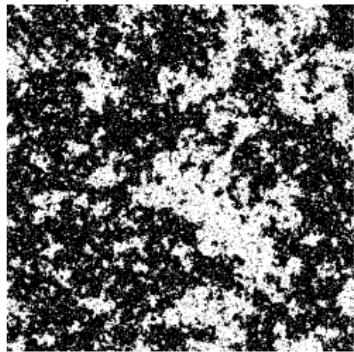


► Dierk Raabe site

# Other CA examples



Sand pile formation



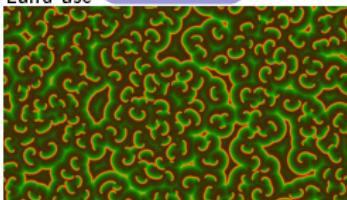
Ising magnetisation

[more info](#)



Land use

[more info](#)



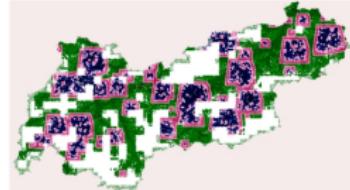
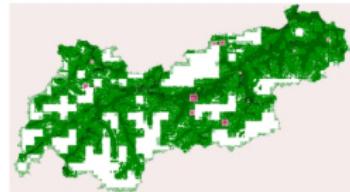
Diffusion

[animation](#)

[info](#)



Fire [more info](#)

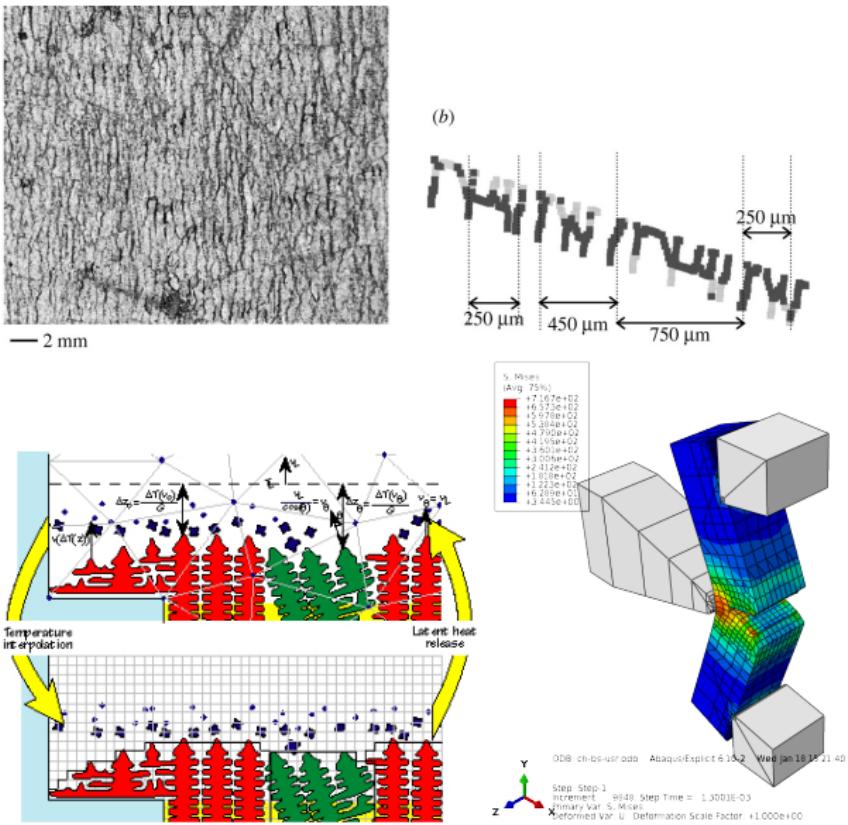


Epidemics, from *The Open Med. Inform.* J. 2(1):70-81, 2008.

[PDF](#)

# More CAFE examples

- ▶ Used for solidification [8], recrystallisation [9] and fracture [10, 11].
- ▶ FE - continuum mechanics - stress, strain, etc.
- ▶ CA - crystals, crystal boundaries, cleavage, grain boundary fracture
- ▶ FE → CA - stress, strain
- ▶ CA → FE - damage variables



## Issues with CrayPAT

71.4%	14,649.9	--	--	USER
<hr/>				
38.7%	7,950.6	913.4	10.3%	cgca_gcupda\$cgca_m3clvg_
24.1%	4,951.2	940.8	16.0%	cgca_clvgp\$cgca_m3clvg_
3.1%	638.0	70.0	9.9%	cgca_pfem_cenc\$cgca_m3pfem_
1.8%	367.5	578.5	61.2%	cgca_hxi\$cgca_m2hx_
1.7%	346.0	196.0	36.2%	cgca_clvgn\$cgca_m3clvg_
<hr/>				

cgca\_gcupda is top in sampling results, but is absent from tracing.  
It is called the same number of times as cgca\_hxi.

29.7%	99.743118	--	--	5,226,813.1  USER
<hr/>				
17.4%	58.326659	36.082315	38.2%	5.0  cgca_clvgp\$cgca_m3clvg_
5.6%	18.876152	5.062089	21.1%	1.0  cgca_pfem_cenc\$cgca_m3pfem_
3.3%	11.145318	15.328335	57.9%	1.0  xx14_
1.7%	5.705317	8.788733	60.6%	5,224,771.1  cgca_clvgn\$cgca_m3clvg_
1.7%	5.689672	1.910819	25.1%	2,035.0  cgca_hxi\$cgca_m2hx_
<hr/>				

# Issues with CrayPAT

All profiling was done with single thread.

```
CrayPat/X: Version 6.2.2 Revision 13378 (xf 13240) 11/20/14 14:32:58
Number of PEs (MPI ranks): 480
```

```
Numbers of PEs per Node: 24 PEs on each of 20 Nodes
```

```
Numbers of Threads per PE: 3
```

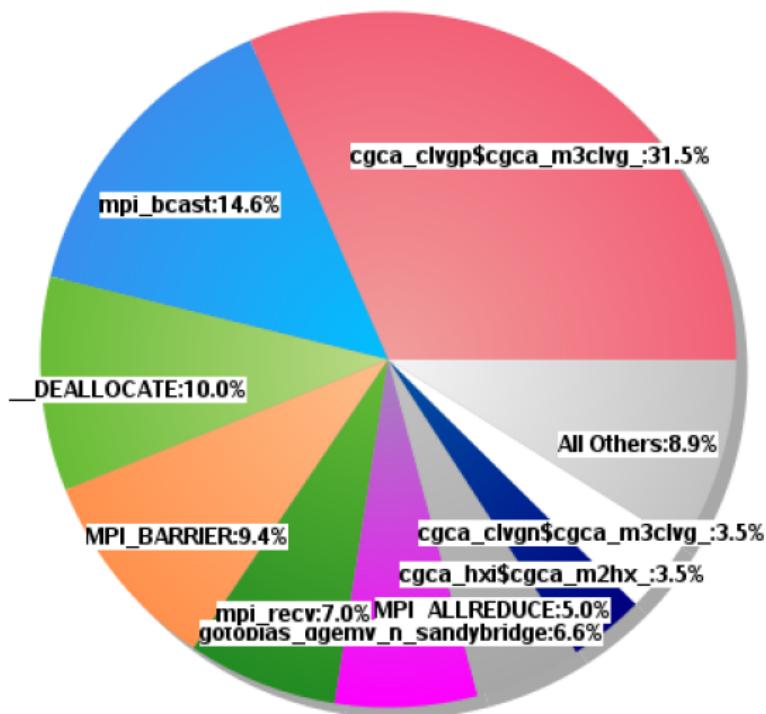
```
Number of Cores per Socket: 12
```

```
Execution start time: Thu Mar 3 13:40:17 2016
```

```
System name and speed: tdsmom 2701 MHz
```

Incorrect number of threads identified by CrayPAT in a tracing experiment of ParaFEM/CGPACK MPI/coarray miniapp with cgca\_gcupda.

# Profiling with cgca\_pfem\_map



Profiling function distribution for ParaFEM/CGPACK MPI/coarray miniapp with cgca\_gcupdn and cgca\_pfem\_map at 7200 cores.

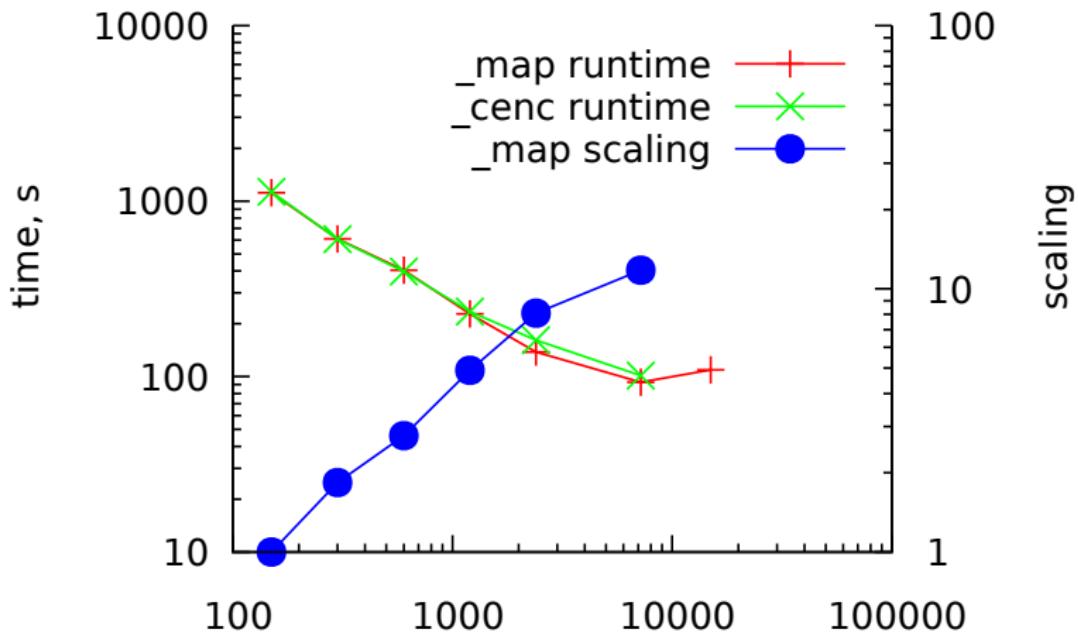
# Profiling with cgca\_pfem\_map

Table 1: Profile by Function

Samp%	Samp	Imb.	Imb.	Group
	Samp	Samp%	Function	
				PE=HIDE
100.0%	9,903.4	--	--	Total
43.6%	4,321.6	--	--	USER
31.4%	3,110.7	589.3	15.9%	cgca_clvgp\$cgca_m3clvg_
3.5%	346.0	513.0	59.7%	cgca_hxi\$cgca_m2hx_
3.5%	342.0	175.0	33.8%	cgca_clvgn\$cgca_m3clvg_
1.2%	116.3	4.7	3.9%	cgca_pfem_map\$cgca_m3pfem_
1.1%	106.8	1,537.2	93.5%	cgca_clvgsd\$cgca_m3clvg_
1.0%	99.9	24.1	19.5%	cgca_sld\$cgca_m3sld_
38.4%	3,803.6	--	--	MPI
14.6%	1,446.6	350.4	19.5%	mpi_bcast
9.4%	932.4	473.6	33.7%	MPI_BARRIER
7.0%	689.5	371.5	35.0%	mpi_recv
4.9%	489.3	76.7	13.6%	MPI_ALLREDUCE
1.5%	145.4	314.6	68.4%	MPI_REDUCE
17.8%	1,766.8	--	--	ETC
9.9%	983.9	8.1	0.8%	_DEALLOCATE
6.6%	652.3	93.7	12.6%	gotoblas_dgemv_n_sandybridge

Raw profiling data  
for ParaFEM/CG-  
PACK MPI/coarray  
miniapp with  
cgca\_gcupdn and  
cgca\_pfem\_map at  
7200 cores.

## Profiling with cgca\_pfem\_map



Runtimes and scaling for ParaFEM/CGPACK MPI/coarray miniapp with `cgca_pfem_map` and `cgca_pfem_cenc`.

`cgca_pfem_map` or `cgca_pfem_cenc` are called only once during the execution of the miniapp. Hence only a minor improvement is obtained, only from about 1000 cores.

- [1] A. Shterenlikht and L. Margetts. Three-dimensional cellular automata modelling of cleavage propagation across crystal boundaries in polycrystalline microstructures. *Proc. Roy. Soc. A*, 471:20150039, 2015.
- [2] A. Shterenlikht, L. Margetts, L. Cebamanos, and D. Henty. Fortran 2008 coarrays. *ACM Fortran Forum*, 34:10–30, 2015.
- [3] A. Shterenlikht. Fortran coarray library for 3D cellular automata microstructure simulation. In M. Weiland, A. Jackson, and N. Johnson, editors, *Proc. 7th PGAS Conf., 3-4 October 2013, Edinburgh, Scotland, UK*, pages 16–24. The University of Edinburgh, 2014.
- [4] D. Henty, A. Jackson, C. Moulinec, and V. Szerem. Performance of Parallel IO on ARCHER, version 1.1. ARCHER White Papers, 2015.
- [5] I. M. Smith, D. V. Griffiths, and L. Margetts. *Programming the Finite Element Method*. Wiley, 5ed, 2014.
- [6] A. Shterenlikht, S. Margetts, L. McDonald, and N. K. Bourne. Towards mechanism-based simulation of impact damage using exascale computing. In *Proc. 19th APS Conf. Shock Compression Condensed Matter SCCM-2015, JUN-2015, Tampa, Florida, USA*, 2015.
- [7] B. Long. Additional parallel features in Fortran. *ACM Fortran Forum*, 35:16–23, 2016.
- [8] C. A. Gandin and M. Rappaz. A coupled finite element-cellular automaton model for the prediction of dendritic grain structures in solidification processes. *Acta Met. and Mat.*, 42(7):2233–2246, 1994.
- [9] C. Zheng and D. Raabe. Interaction between recrystallization and phase transformation during intercritical annealing in a cold-rolled dual-phase steel: A cellular automaton model. *Acta Materialia*, 61:5504–5517, 2013.

- [10] A. Shterenlikht and I. C. Howard. The CAFE model of fracture – application to a TMCR steel. *Fatigue Fract. Eng. Mater. Struct.*, 29:770–787, 2006.
- [11] S. Das, A. Shterenlikht, I. C. Howard, and E. J. Palmiere. A general method for coupling microstructural response with structural performance. *Proc. Roy. Soc. A*, 462:2085–2096, 2006.